

Online Algorithms

Offline algorithms:

- All inputs are given in advance.
- E.g., given n numbers, the merge-sort algorithm can sort the numbers in $O(n \log n)$ time.


Online algorithms:

- Inputs (requests) are given **one** at a time;
 - the algorithm must serve each input before it receives the next one.
- Usually assume no knowledge of the future inputs.
- Why difficult?
 - the offline version is NP-complete; and/or
 - lack of future information.

Example: Ski Rental

- You are a novice skier and have no idea whether you really like skiing (and when you will give up skiing).
- Each time you go to ski, you need to make a decision of renting or buying a pair of skis.
- Rental: \$100; Buying cost: \$1200
- What is a good strategy?
 - (a) Buy the skis the first time you go to ski.
 - (b) Always rent the skis
 - (c) Buy the skis after you've skied 3/6/12 times.

Example: Paging algorithms



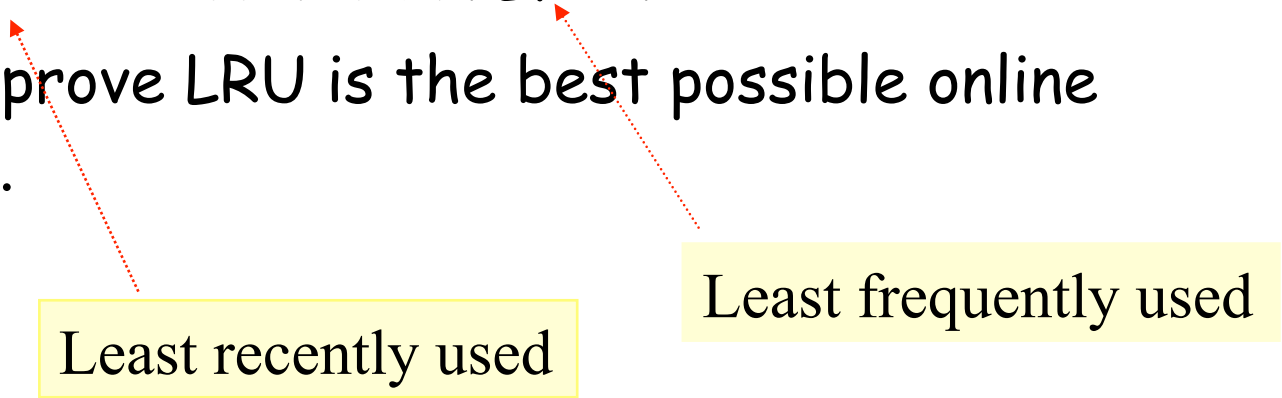
- A memory request is the range $[1, N]$.
- To serve a request, bring the requested memory item into the cache if it is not already there (i.e., a miss).
- A paging algorithm (like LRU, LFU) determines which item in the cache to be evicted so as to make room for the requested item.

Why studying online algorithms?

- For many real-life online problems, we do have some algorithms (heuristics, ideas) for solving them.
 - In many cases, we believe some algorithms are better than others, but often without justification.
- Theory folks always want to explain (or disprove) such phenomena.
- A typical example is the paging problem.

Can we prove LRU is better than LFU ?

Or even better, prove LRU is the best possible online paging algorithm.



Least recently used

Least frequently used

Performance measure

- Cost = number of memory access
- Assumption: An algorithm assesses the memory only when there is a miss.
 - Given an algorithm that may access the memory at any time, we can construct another algorithm that accesses the memory only when a miss occurs with no increase in total number of access.
- Cost: Serving a request is associated with a cost.
E.g., paging algorithms: hit \rightarrow 0; miss \rightarrow 1.

Competitive ratio

- Cost: Serving a request is associated with a cost.
 - E.g., paging algorithms: hit $\rightarrow 0$; miss $\rightarrow 1$.
- Given a sequence I of requests, we want an online algorithm A to minimize the total cost.
- Let $A(I)$ denote the total cost that A needs to serve I .
- We say that
 - A is (asymptotic) c -competitive: $\forall I, A(I) \leq c \text{ OPT}(I) + d$, where $\text{OPT}(I)$ denotes the smallest possible cost to serve I , and d is a constant;

P.S. Roughly speaking, 5-competitive implies a cost at most 5 times of the minimum cost in the worst case.

 - A is optimal: $\forall I, A(I) = \text{OPT}(I)$.

c is called the **competitive ratio** or the competitiveness coefficient.

Why competitive analysis ?

- Suppose we simply measure the worst-case cost in terms of the length (denoted n) of input sequence. Say, $\log n$, n , etc.
- Meaningless in many cases. For example, we can show that given **any** paging algorithm **A**, there exists a sequence I such that $A(I) = \text{the length of } I = n$.
- What does it mean? **All online paging algorithms are equally bad.**
- Intuitively, an algorithm shouldn't be regarded as poor if whenever it performs poorly, the optimal offline algorithm also performs poorly.

LRU is better than LFU

- Lower bound: No paging algorithm has competitive ratio smaller than k , where k is the size of the cache.

[Sleator and Tarjan 1985]

- LRU is k -competitive.

[Sleator and Tarjan 1985; Goemans 1993]

- LFU cannot achieve a bounded competitive ratio. (That means, there are some input sequences for which LFU performs poorly.)

Outline

1. LRU is k -competitive.
2. LFU cannot achieve a constant competitive ratio.
3. Lower bound: no (deterministic) online algorithm can be better than k -competitive
4. Randomized algorithm: $O(\ln k)$ -competitive

LRU is k -competitive

LRU: When eviction is necessary, pick the item whose "recent use" was earliest.

Let I be any sequence of jobs.

Divide I into phases as follows:

Phase 1: k distinct items	Phase 2 : k distinct items	Phase 3 : k distinct items
-----------------------------	------------------------------	------------------------------

Phase 1: the **longest** sequence requesting for at most k **distinct** items.

....

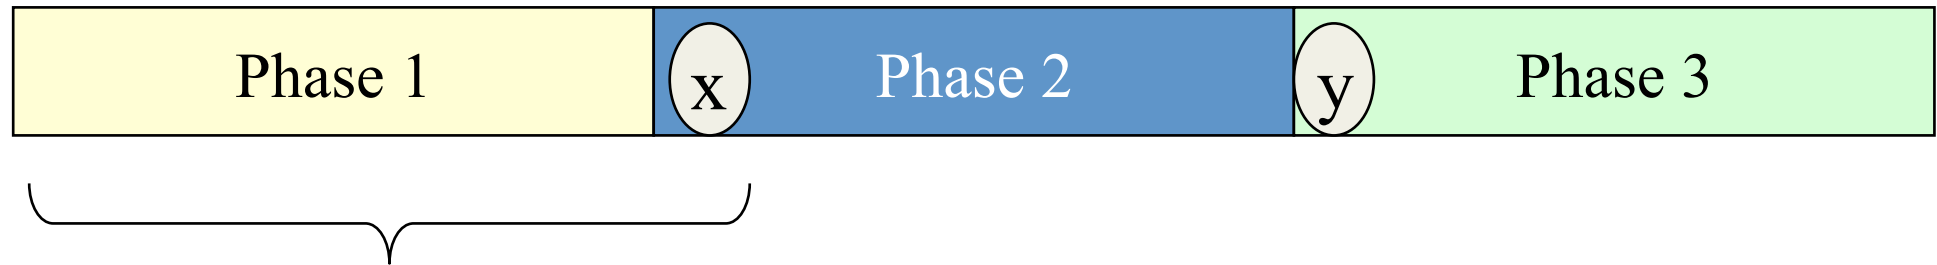
Phase i : the **longest** sequence following Phase $i-1$ requesting at most k **distinct** items.

....

Proof Framework

- Optimal offline algorithm: at least one miss in every phase;
cost ≥ 1 .
- LRU: at most k misses in every phase; cost $\leq k$.

Opt = optimal offline algorithm



Let us consider the subsequence S comprising Phase 1 and the first request of Phase 2.

By definition of a phase, S requests $k+1$ distinct items.

Opt must incur a miss. Why?

Suppose on the contrary that there isn't a miss.

Then Opt with a cache of k items can serve requests for $k+1$ distinct items.


A contradiction occurs.



Let us consider the subsequence S comprising Phase 2 (except the first item x) and the first request of Phase 3.

S contains k distinct items $\neq x$.

Opt must incur a miss. Why?



Suppose on the contrary that there isn't a miss.

Then Opt with a cache containing x and $k-1$ items can serve requests for k distinct items not equal to x .


A contradiction occurs.



Let us consider the subsequence S comprising Phase i (except the first item z) and the first request of Phase $i+1$.

S contains k distinct items not equal to z .

Opt must incur a miss.





Conclusion on opt

Except the last phase, we can charge at least one miss to every phase.

Total cost \geq number of phases - 1

LRU causes at most k misses within a phase.

Lemma 1: Within a phase, once an item p is brought into the cache, p won't be evicted by LRU.




Corollary: Within a phase, each of the k distinct items can cause at most one miss, which can only occur at its first occurrence.

In other words, the number of misses is at most k .

Proof of Lemma 1 – by contradiction

Suppose on the contrary that within a phase, an item p is evicted after it has been requested.



Assume this eviction occurs at the request I_y .
Denote I_x be the last request for p before I_y . I.e., $x < y$.

- **Immediately after** I_x : p is the most recently used item.
- **Just before** I_y : p is the least recently used (i.e., k -th most recently used) item.
- To make p the least recently used item and then evicts it, how many distinct items other than p must be requested from I_{x+1} to I_y ?

From I_{x+1} to I_y

At the point the 1st item other than p is requested, p is still the most recently used.

At the point the 2nd distinct item other than p is requested, p is the 2nd-most recently used.

At the point the 3rd distinct item...., p is the 3rd most recently used.

...

At the point the $(k-1)^{\text{st}}$, p is the $(k-1)^{\text{st}}$ most recently used.

Thus, if eviction is required here, p is not yet the candidate according LRU.

In other words, we need at least k distinct items other than p to cause an eviction of p .

This contradicts that a phase contains k distinct items.

LRU is k -competitive

Consider any sequence I .

Suppose I is partitioned into m phases.

$$\text{LRU}(I) \leq km$$

$$\text{Opt}(I) \geq m - 1$$

$$\text{Therefore, } \text{LRU}(I) \leq km \leq k(\text{Opt}(I) + 1) \leq k \text{Opt}(I) + k$$

Outline

1. LRU is k -competitive.
2. LFU cannot achieve a constant competitive ratio.
3. Lower bound: no (deterministic) online algorithm can be better than k -competitive
4. Randomized algorithm: $O(\ln k)$ -competitive

LFU is not competitive

LFU: When eviction is necessary, pick the item p in the cache that has the smallest number of access so far.

Lemma. For any positive integer c , there exists I such that $\text{LFU}(I) \geq c \text{Opt}(I)$.

LFU is not competitive

LFU: When eviction is necessary, pick the item p in the cache that has the smallest number of access so far.

Lemma. For any positive integer c , there exists I such that $\text{LFU}(I) \geq c \text{Opt}(I)$.

$I: (p_1)^c (p_2)^c (p_3)^c \dots (p_{k-1})^c (p_k p_{k+1})^{c-1}$

Opt: 1 miss

LFU is not competitive

LFU: When eviction is necessary, pick the item p in the cache that has the smallest number of access so far.

Lemma. For any positive integer c , there exists I such that $\text{LFU}(I) \geq c \text{Opt}(I)$.

$I: (p_1)^c (p_2)^c (p_3)^c \dots (p_{k-1})^c (p_k p_{k+1})^{c-1}$

$\text{Opt}: 1 \text{ miss}$

$\text{LFU}: 2c - 3 \text{ miss}$

$(p_1)^c (p_2)^c (p_3)^c \dots (p_{k-1})^c p_k p_{k+1} p_k p_{k+1} p_k p_{k+1} \dots p_k p_{k+1}$

Outline

1. LRU is k -competitive.
2. LFU cannot achieve a constant competitive ratio.
3. **Lower bound**: no (deterministic) online algorithm can be better than k -competitive
4. Randomized algorithm: $O(\ln k)$ -competitive

Lower bound: Competitive ratio $\geq k$

Let A be any online paging algorithm.

Claim: There exists a sequence I of n requests such that
 $A(I) = |I| = n$.

Proof: I comprises memory requests the range $[1..k+1]$.

At any time the cache contains at most k items.

The next request is simply the item not in the cache.

Therefore, $A(I) = n$.

Framework

Let A be any online paging algorithm.

Fact. There exists a sequence \mathbf{I} of n requests such that $A(\mathbf{I}) = n$.

How to argue that competitive ratio of $A \geq k$?

We need to show that the optimal offline algorithm O when working on \mathbf{I} incurs a cost $\leq n/k$.


Then competitive ratio of $A \geq A(\mathbf{I})/\text{Opt}(\mathbf{I}) \geq n / (n/k) = k$.

Offline algorithm: $Opt(I) \leq LNU(I)$

Consider the following offline algorithm LNU (latest next use; LFD):


Initially, the cache is loaded with any k items.

When eviction is necessary, choose the item in the cache that will not be used for the longest period of time.



- LNU is an offline algorithm. $Opt(I) \leq LNU(I)$ for all I .


Lemma. Consider any request sequence I in the range $[1, k+1]$.
Between every two misses of LNU, there must be at least $k-1$ hits.



Therefore, one miss is followed by at least $k-1$ hits before any miss, and $LNU(I) \leq \lceil n / k \rceil$.

Corollary: For any sequence I of memory requests in the range $[1, k+1]$, $LNU(I) \leq \lceil n / k \rceil$

A miss of LNU




With respect to a request sequence I , suppose LNU incurs a miss at the x -th request ($I_x = p'$) and evicts item p from the cache.

Let Q be the set of the other $k-1$ items currently in the cache.

Immediately after I_x , what request will cause a hit?

What request will cause a miss?

A miss of LNU



With respect to a request sequence I , suppose LNU incurs a miss at the x -th request ($I_x = p'$) and evicts item p from the cache.

Let Q be the set of the other $k-1$ items currently in the cache.

Immediately after I_x , what request will cause a hit?

$$Q \cup \{p'\} = \{1, \dots, k+1\} - \{p\}$$

What request will cause a miss? p



Two cases to consider:

- After I_x , p is not requested again and LNU will incur a zero cost.
- After I_x , p is requested at I_y .



Why item p is chosen by LNU for eviction when serving I_x ?

Because for each item q in Q , q will be requested before p .

In other words, **after** I_x and **before** I_y , each of the $(k-1)$ items in Q is requested at least once, and there are at least $k-1$ hits.

Thus, a miss must be followed by at least $k-1$ hits,
and $LNU(I) \leq n / k$.

Conclusion

- For any algorithm A , there exists a sequence I of n requests such that $A(I) = n$.
- $\text{Opt}(I) \leq \text{LNU}(I) \leq n / k$.
- The competitive ratio of $A \geq A(I) / \text{Opt}(I) \geq k$.

Outline

1. LRU is k -competitive.
2. LFU cannot achieve a constant competitive ratio.
3. Lower bound: no (deterministic) online algorithm can be better than k -competitive
4. Randomized algorithm: $O(\ln k)$ -competitive

Can we do better?

Randomized online algorithms: In the worst case (input), can the **expected** value of the competitive ratio $< k$?

- Evict an item chosen randomly: k -competitive
[Raghavan & Snir 1989]
- A random marker algorithm: $2 H_k$ -competitive, where $H_k = 1 + 1/2 + 1/3 + \dots + 1/k \leq 1 + \ln k$
[Fiat, Karp, Luby, McGeoch, Sleator, Young 1991]
[Achlioptas, Chrobak, and Noga 1996]
- Lower bound: No randomized paging algorithm has competitive ratio smaller than H_k .

Restricting the inputs to model the locality of reference:

- LRU is better than FIFO ...

[Borodin, Irani, Raghavan, Schieber 95] [Chrobak & Noga 98]...

Algorithm Mark

Initially, all items in the cache, if any, are considered to be unmarked.

When an item in the cache is requested, **mark** that item.


Whenever eviction is necessary, choose an **unmarked** item in the cache randomly.

- The new item to be brought into the cache will be **marked**.
- If all items are **marked**, **unmark** ALL of them first.

Remark: Most of the time, recently used pages are given priority to stay in the cache. But occasionally, ignore the history.

Analysis

For any input sequence I , partition I into phases as before (precisely, Phase i captures the **longest** sequence following Phase $i-1$ involving at most **k distinct** items).



The k distinct items requested in Phase 1 must all be marked and kept in the cache until the end of Phase 1.

Unmark
all.

The first item of Phase 2 would trigger unmark-all and then an eviction.

At Phase i

Phase $i-1$: k distinct items

Phase i : k distinct items



The k distinct items requested in Phase $i-1$ must be in the cache of Mark.
Call these items old items to Phase i .

Phase i requests for k distinct items. Suppose there are n_i new (not old) items and $k - n_i$ old items.

NB. $n_1 = k$; for $i > 1$, $n_i \geq 1$

Phase i-1: k distinct items

Phase i: k distinct items

Phase i requests n_i new items and $k - n_i$ old items.

How many misses for **Opt** (optimal offline algorithm)? At least ..

In Phases i-1 and i, there are $k + n_i$ distinct items requested.

Number of miss $\geq n_i$

Opt' s total cost $\geq n_2 + n_4 + n_6 + \dots$

Alternatively, total cost $\geq n_1 + n_3 + n_5 + \dots$

In conclusion, Opt' s total cost $\geq \frac{1}{2} [n_1 + n_2 + n_3 + n_4 + \dots] = \frac{1}{2} \sum n_i$

Phase $i-1$: k distinct items

Phase i : k distinct items

Phase i requests n_i new items and $k - n_i$ old items.

How many misses for **Mark** in Phase i ? At most ...

Misses due to new items in Phase i : exactly n_i

In the worst case, new items should be requested before old items (so that each old item has a better chance to cause a miss).

Miss due to old items: The exact number depends on the random choice of items for eviction (to make room for the new items).

Phase i-1: k distinct items

Phase i: k distinct items



Phase i requests n_i new items and $k - n_i$ old items.

	Probability in cache	Miss probability
1 st old item requested	$(k - n_i) / k$	n_i / k
2 nd old item requested	$(k - n_i - 1) / (k - 1)$	$n_i / (k - 1)$
.....		
$(k - n_i)$ -th old item	$1 / (n_i + 1)$	$n_i / (n_i + 1)$

Remark. When the j -th old item is requested,

- all the n_i new items and the first $(j - 1)$ old items must be in the cache;
- the j -th old item as well as every other old item have the same probability to be in the cache. I.e., $(k - (j - 1))$ items fight for $(k - n_i - (j - 1))$ positions.

Expected # of miss in phase i due to old items:

$$1 \times (n_i / k) + 1 \times (n_i / k - 1) + \dots + 1 \times (n_i / n_i + 1)$$

$$= n_i [1 / k + 1 / k - 1 + \dots + 1 / n_i + 1]$$

$$= n_i [H_k - H_{ni}]$$

$$\text{Expected \# of miss in phase i} = n_i + n_i [H_k - H_{ni}] \leq n_i H_k$$

$$\text{Expected \# of miss in all phases} = \sum n_i H_k = H_k \sum n_i$$

Conclusion

Opt' s total cost: $\geq \frac{1}{2} \sum n_i$

Mark (expected cost) : $\leq H_k \sum n_i$

Competitive ratio (expected value): $2 H_k$

Reference

- Online computation & competitive analysis, Cambridge, 1998, Borodin & El-Yaniv