

Compressed text indexing

- **FM-index**: compressed Burrows Wheeler Transform (**BWT**) [Ferragina, Manzinni, focs 00]
- **Compressed suffix arrays**: [Grossi & Vitter, stoc 00]

Space: $O(n)$ **bits** (instead of $O(n)$ words)

Today's lecture

- BWT for pattern matching
- $2n$ bits for DNA; i.e, same size as the text

Definition of BWT [Burrows & Wheeler 94]




- $T[0..n]$ -- a text of $n+1$ characters, where $T[n] = \$$ (a special character).
- $SA[0..n]$ -- $SA[i] = j$ if $T[j..n]$ is the i^{th} lexicographically smallest suffix.

The rank of $T[j..n]$ is said to be i .

- $BWT[i] = T[j-1]$, where $j = SA[i]$.

I.e., $BWT[i]$ is the character immediately before the i^{th} smallest suffix.



Example

- $T = ACACGT\$$
- $BWT = T\$CAACG$
- $BWT[i]$ is the character immediately before the i^{th} suffix.

	SA	BW T	Sorted Suffixes
0	6	T	\$
1	0	\$	ACACGT\$
2	2	C	ACGT\$
3	1	A	CACGT\$
4	3	A	CGT\$
5	4	C	GT\$
6	5	G	T\$


Backward searching

BWT (plus some auxiliary function) allows us to search for a pattern in a backward manner.

Given a pattern $P[1..m]$, compute

- the (smallest) rank of $P[m]$ w.r.t. the suffix array of T
- the (smallest) rank of $p[m-1..m]$
- ...
- the (smallest) rank of $p[2..m]$
- the (smallest) rank of $p[1..m]$

NB. The largest rank is computed in a similar way.



Count(x)

- For any character x , let $\text{Count}(x)$ be the total # of characters in the text T that are smaller than x .

E.g., if $T = \text{ACACGT\$}$, then


$\text{Count}(A) = 1;$

$\text{Count}(C) = 3;$

$\text{Count}(G) = 5$

$\text{Count}(P[m])$


- (smallest) Rank of $P[m..m] = \text{Count}(P[m])$
E.g., $P = \text{ACG}; \text{Count}(G) = 5$



From $P[j+1..m]$ to $P[j..m]$

Let $P[j+1..m] = w$, and let $P[j] = x$.

Then $P[j..m] = \textcolor{red}{x}w$.



Suppose we've already computed the (smallest) rank of w , which is equal to i .

Next, we want to compute the (smallest) rank of $\textcolor{red}{x}w$ (denoted i' in the figure).

0	
1	
2	
.	
.	
	X...
	X...
i'	XW..
	XW..
	X...
i	W..
	W..
	W..

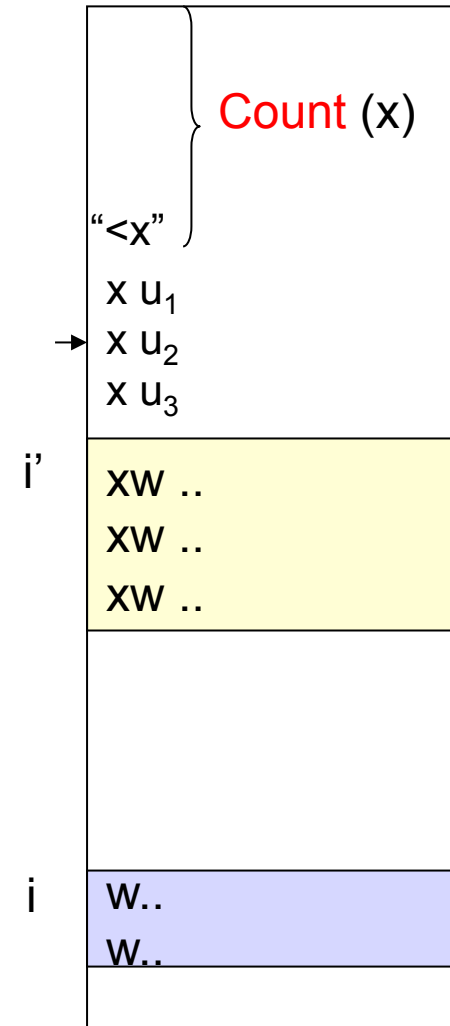
Computing rank of xw

Rank of xw = $\text{Count}(x) +$
 $\#$ of suffixes xu such that $u <_l w$

Consider a suffix xu with $u <_l w$.

- $u <_l w \Leftrightarrow \text{rank of } u < (\text{smallest}) \text{ rank of } w$.
- Recall that $(\text{smallest}) \text{ rank of } w = i$.

Let rank of $u = k$. Then $k < i$.



Computing rank of xw

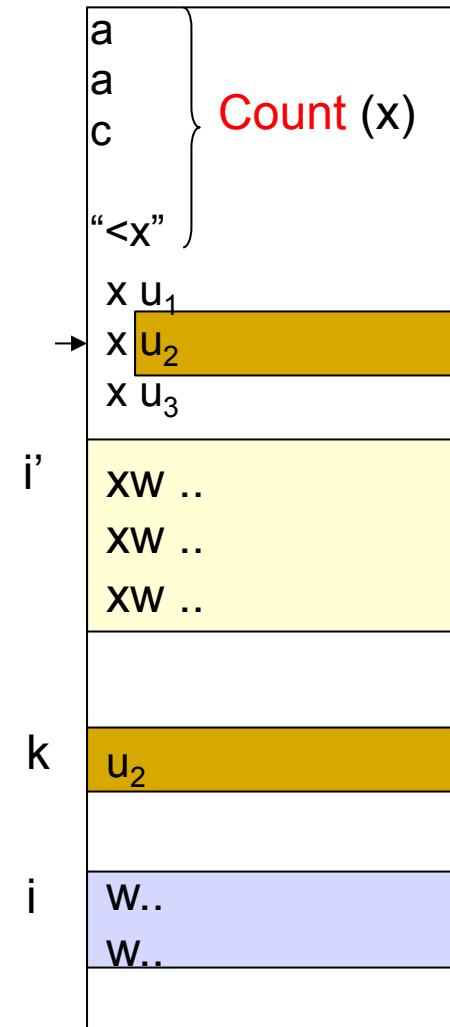
Rank of $xw = \text{Count}(x) +$
 $\#$ of suffixes xu such that $u < w$

Consider a suffix xu of T with $u < w$.

- $u < w \Leftrightarrow \text{rank of } u < (\text{smallest}) \text{ rank of } w$.
- Recall that (smallest) rank of $w = i$.
 Let rank of $u = k$. Then $k < i$.



- Then $k < i$ and $BWT(k) = "x"$



Use BWT to compute rank of xw

Rank of $xw = \text{Count}(x) +$
 $\#$ of suffixes xu such that $u <_l w$

Each suffix xu with $u <_l w$
 corresponds uniquely to
 an index $k < i$ such that $\text{BWT}(k) = "x"$.

$\#$ of suffixes xu of T such that $u < w$
 $= \#$ of indices $k < i$ such that $\text{BWT}(k) = "x"$
 $= \#$ of x 's in $\text{BWT}[0..i-1]$

	a	}	Count (x)
	a		
	c		
	"<x"		
	x u ₁		
→	x u ₂		
	x u ₃		
i'	xw ..		
	xw ..		
	xw ..		
k	u ₂		
i	w..		
	w..		

Example

- $T = ACACGT\$$
- $P = ACG$
- Notation: **Appear** $[i,x] =$
of x in $BWT[0..i-1]$
- Rank of $G = \text{Count}(G) = 5$

	SA	BWT	Sorted Suffices
0	6	T	\$
1	0	\$	ACACGT\$
2	2	C	ACGT\$
3	1	A	CACGT\$
4	3	A	CGT\$
5	4	C	G T\$
6	5	G	T\$

Example

- $T = ACACGT\$$
- $P = \textcolor{red}{ACG}$
- Notation: $\text{Appear}[i,x] = \# \text{ of } x \text{ in } BWT[0..i-1]$
- Rank of $\textcolor{red}{G} = \text{Count}(G) = 5$
- Rank of $\textcolor{red}{CG} = \text{Count}(C) + \text{Appear}(5,C) = 3 + 1 = 4$

	SA	BWT	Sorted Suffices
0	6	T	\$
1	0	\$	ACACGT\$
2	2	C	ACGT\$
3	1	A	CACGT\$
4	3	A	$\textcolor{red}{CG}$ T\$
5	4	C	GT\$
6	5	G	T\$

Example

- $T = ACACGT\$$
- $P = \textcolor{red}{ACG}$
- Notation: $\text{Appear}[i,x] = \# \text{ of } x \text{ in } BWT[0..i-1]$

Rank of $\textcolor{red}{G} = \text{Count}(G) = 5$

Rank of $\textcolor{red}{CG} = \text{Count}(C) + \text{Appear}(5,C) = 3 + 1 = 4$

Rank of $\textcolor{red}{ACG} = \text{Count}(A) + \text{Appear}(4,A) = 1 + 1 = 2$

	SA	BWT	Sorted Suffices
0	6	T	\$
1	0	\$	ACACGT\$
2	2	C	$\textcolor{red}{ACG}$ T\$
3	1	A	CACGT\$
4	3	A	CGT\$
5	4	C	GT\$
6	5	G	T\$

Searching details

For any pattern $P[1..m]$, for any $1 \leq j \leq m$, let

- $first_j$ be the SA index (rank) of the first suffix matching $P[j..m]$; and
- $last_j$ be the SA index (rank) of the last suffix matching with $P[j..m]$.
 - $first_m = \text{Count}(P[m])$
 - $last_m = \text{Count}(P[m] + 1) - 1$
 - ...
 - $first_j = \text{Count}(P[j]) + \text{Appear}(P[j], first_{j+1})$
 - $last_j = \text{Count}(P[j]) + \text{Appear}(P[j], last_{j+1} + 1)$

Preprocessing


- To save time, we perform preprocessing on **Appear & Count**.

$o(|\Sigma|n)$ bits are sufficient to allow very efficient retrieval of **Count** and **Count**.

- FM-index is a compressed version of BWT (e.g. using move-to-front and run-length encoding). However, for DNA ($|\Sigma| = 4$), the compression is not effective and doesn't save much space.
- Finding the actual occurrences requires additional auxiliary data structures (SA sampling).

Appear: simple implementation

- BWT itself: $n \log |\Sigma|$ bits
- Auxiliary data structure for computing **Appear**: $o(n|\Sigma|)$ bits
- Take $L = O(\log n)$, for every L^2 BWT characters, we store the **Appear**() value.
- For every L BWT character, we store the difference w.r.t. the last L^2 BWT position. This takes only $\log(L^2)$ bits as the value must be within L^2 .



- In total, this takes $|\Sigma| (n / L^2 (\log n) + n / L (\log L^2)) = O(|\Sigma| (n / \log n + n \log \log n / \log n)) = o(n|\Sigma|)$ bits.

BWT segments of $\log n$ characters

- Given a segment of $\log n$ characters and a character x , we want to find how many times x appears in the segment using $O(1)$ time.
- Each character is represented by $\log \Sigma$ bits.
- Consider a short segment of $h=10$ characters.
 - Index: In advance count the number of x in every possible segment of h characters and store the counts in a table (array) A indexed by " h characters" (i.e., a binary sequence of $h \log \Sigma$ bits).
 - Counting: Given any h characters, counting takes $O(1)$ time.
- Let $h = \sim \frac{1}{2} \log n / \log \Sigma$.

For each character x , build a table A_x

Table A_x : $2^{\frac{1}{2} \log n} = n^{\frac{1}{2}}$ entries

Alternative: Rank & Select

- For each character x , define a bit vector V such $V[i] = 1$ if and only if $BWT[i] = x$.
- Build a rank-and-select data structure for V , then $Appear[x,i]$ can be computed using a rank operation.
- Space: for each character, $o(n)$ bits (precisely, $(1 + o(1)) n \log \log n / \log n + O(n / \log n)$ bits).

Space-efficient Construction of BWT

- Don't build SA first. Instead build BWT direct.
- [Hon et al focs 03 & Algorithmica 07, Lippert et al. JCB 05]
- Requires only $2.5n$ bits of memory; takes 20 minutes to build a BWT for the human genome (on an ordinary PC).
- Converting BWT to FM-index or CSA is even faster.
- For DNA, BWT (uncompressed FM-index) has the best performance regarding space & time.


Survey paper

- Compressed full-text indexes, Navarro & Makinen, ACM Computing Surveys, April 2007

Text Indexing

Build an index for a DNA sequence T (say, human genome) in the main memory to facilitate

- Exact pattern matching
- Approximate string matching (semi-global alignment)



Approximate string matching

Not trivial even when the number of errors, k , is small.

- Suffix tree: $O(n)$ words, $O(|P|^k + \text{occ})$ time
- $O(n \log^k n)$ -word index: $O(|P| + \log^k n \log \log n + \text{occ})$ time [Cole et al. STOC 04]
- $O(n \log^{k-1} n)$ -word index, same k -error matching time [Chan et al. ESA 06]
- $O(n)$ -word index: $O(|P| + \log^{k(k+1)} n \log \log n + \text{occ})$ time [Chan et al. CPM 06]

These indexes occupy more space than a suffix tree. Too big for indexing long DNA sequences.