



# Job scheduling

- Given a set of jobs, find a schedule to run the jobs on one or more processors.
- Jobs have different release times, sizes (amount of work in terms of processing time on a processor), and perhaps deadlines.
- A processor can only process a job at a time. Preemption is allowed.



# Scheduling objectives

- Offline scheduling: all jobs are known in advance.  
Online scheduling: jobs are each known only at their release time
  - requires continuous adjustment of the schedule as more jobs are known.
- For jobs without deadlines:
  - Minimize average/total response time (flow time)
  - Minimize average slowdown (stretch).
- For jobs with deadlines:
  - Hard deadlines: all jobs must be completed before the deadlines (e.g., nuclear reactor)
  - Soft (firm) deadlines: some jobs can be missed; maximize the total work of the jobs that can be completed by their deadlines.

# Online flow (response) time scheduling

- There is a pool of  $m \geq 1$  identical processors.
- Jobs are released at arbitrary times.
- Each job can be executed on only one processor at a time.
- The work (processing time) of a job is known at its release time.; no deadline.
- An online scheduler, based on the jobs released so far, adjusts the (future) ordering of job processing.
- Flow time of a job = completion time - release time.
- Objective: Complete all jobs with total flow time minimized.

J6

A new job may arrive  
at any time.

Online scheduler

NB. Jobs may have  
different processing  
times.

J3

J5

J2

Processor 1

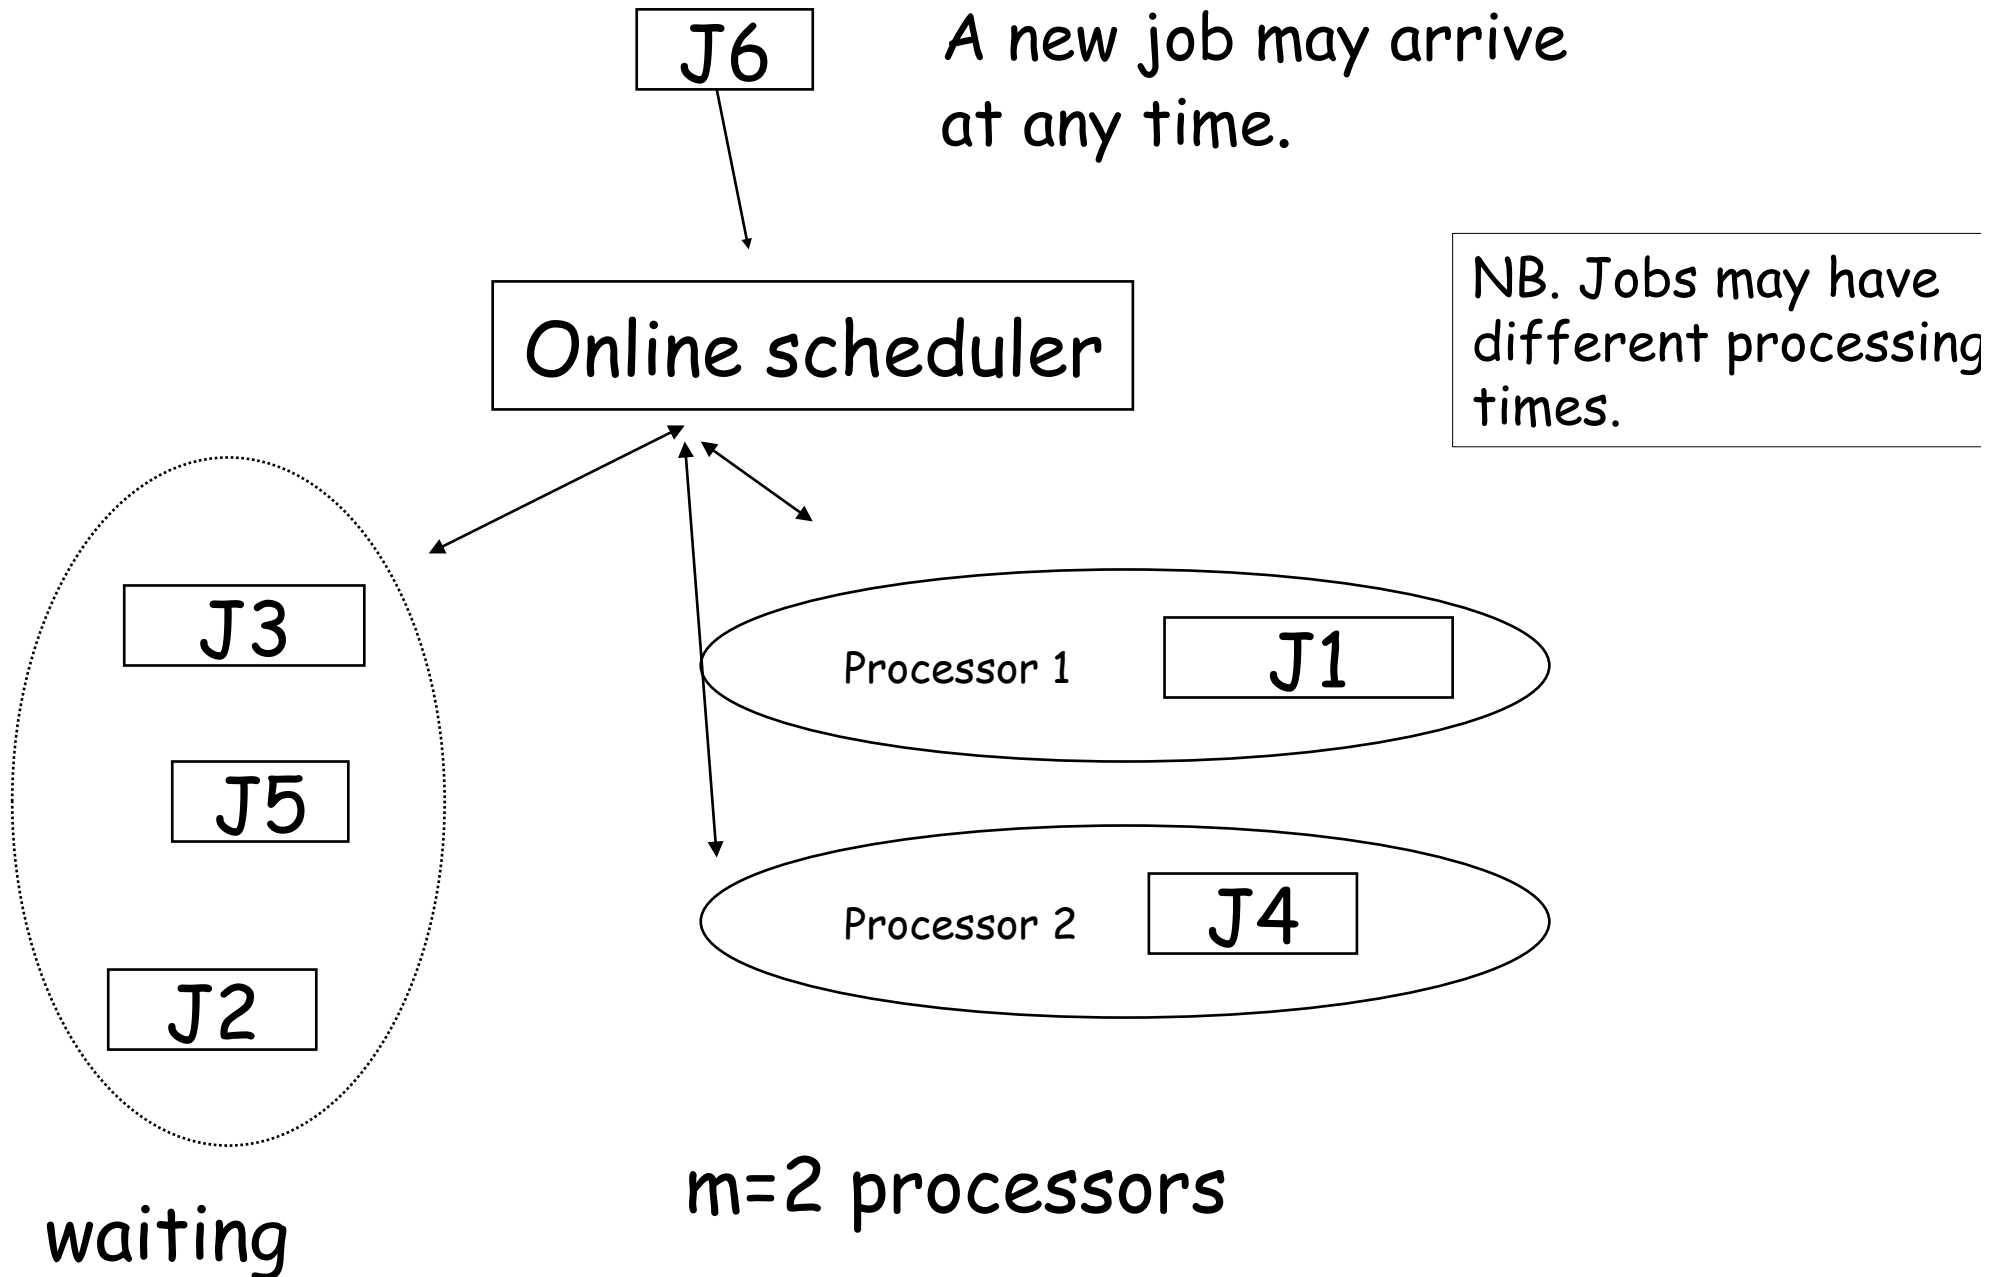
J1

Processor 2

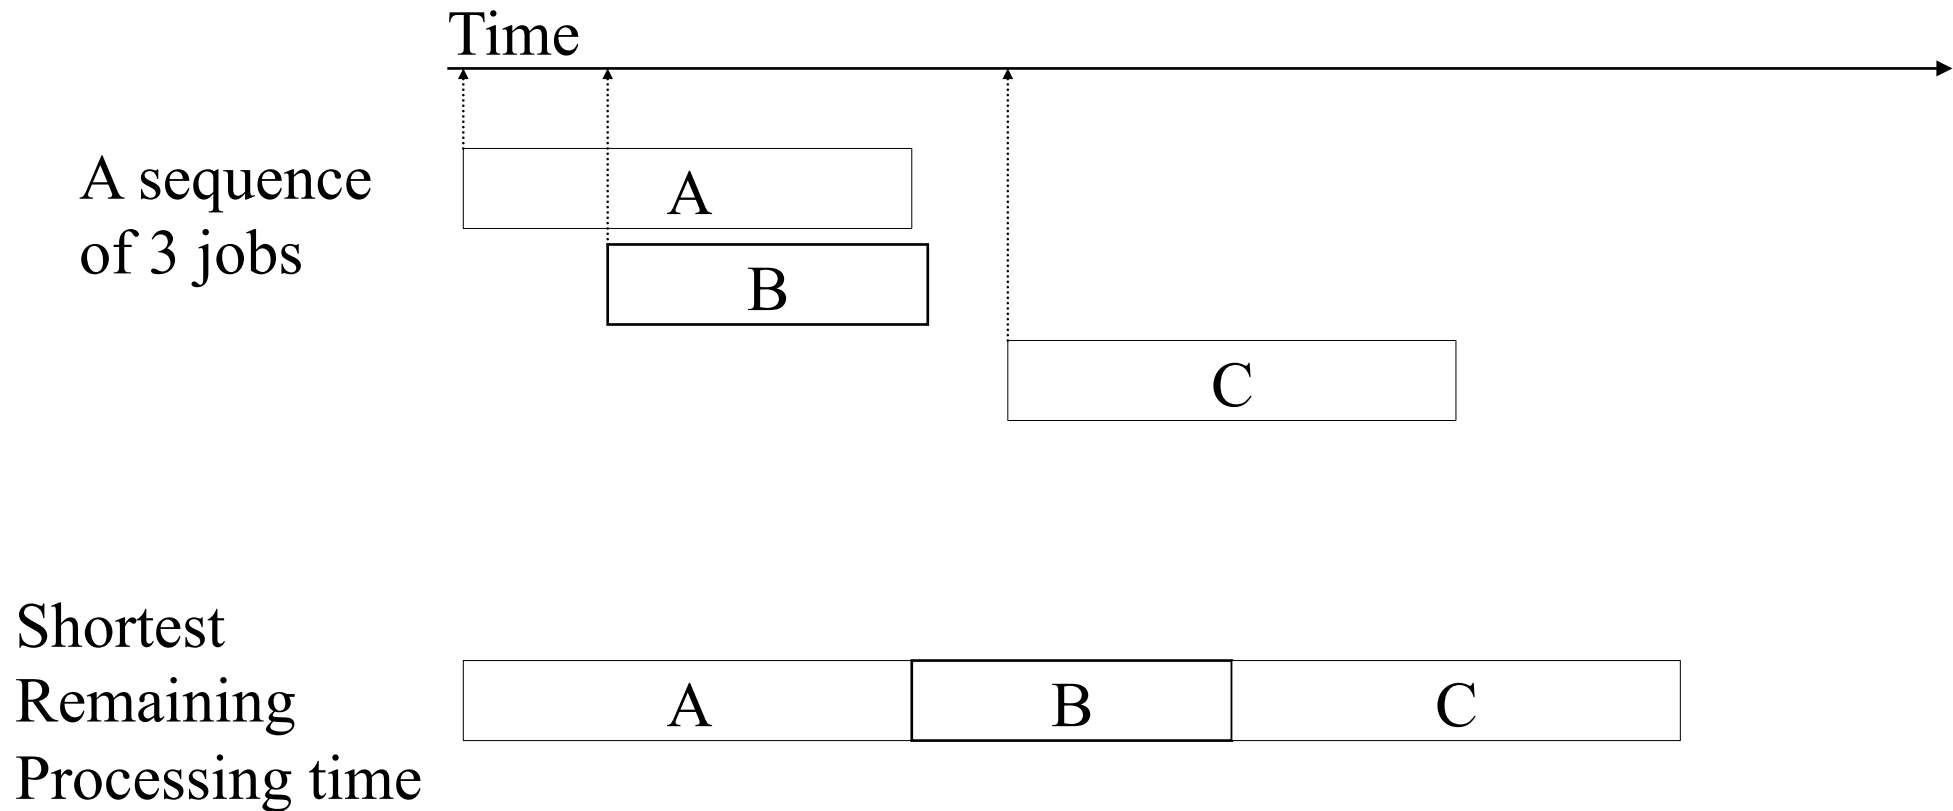
J4

waiting

$m=2$  processors



# Example: $m = 1$ processors



# One processor: SRPT minimizes flow time

SRPT: At any time, schedule the job with the shortest remaining work (processing time).

Theorem. For one-processor scheduling, SRPT is 1-competitive for flow time.


# Proof: a transformation argument

Given a job sequence  $I$ , let  $S$  be a schedule that achieves the minimum total flow time, and that doesn't follow SRPT.

Let  $t$  be the first time  $S$  schedules a job  $J_1$  that is not the job with the shortest remaining work.

Let  $J$  be the job with the shortest remaining work at  $t$ .

Consider the schedule of  $J_1$  &  $J$  in  $S$ .




# Proof: a transformation argument

Given a job sequence  $I$ , let  $S$  be a schedule that achieves the minimum total flow time.


Let  $t$  be the first time  $S$  schedules a job  $J_1$  that is not the job with the shortest remaining work.

Let  $J$  be the job with the shortest remaining work at  $t$ .


Consider the schedule of  $J_1$  &  $J$  in  $S$ .



Rearrange  $S$  so that  $J$  is processed before  $J_1$  after time  $t$ .







In S, let  $x$  be the time S finishes one of J and J<sub>1</sub>. Let  $y$  be the time S finishes both.

- (J<sub>1</sub> & J)'s flow time = total completion time - total release time  
 $= x + y - \text{release}(J_1) - \text{release}(J)$

In S', J will be completed before J<sub>1</sub>.

- J must be completed strictly before  $x$ .
- J<sub>1</sub> will be completed at  $y$ .
- (J<sub>1</sub> & J)'s flow time  $< x + y - \text{release}(J_1) - \text{release}(J)$

This contradicts that S minimizes the total flow time.

# multi-processor flow time scheduling (m = 1)

- SRPT is no longer optimal (1-competitive).
- Example: 2004 jobs,  $m = 2$  processors
  - At time 0 : four jobs with 10, 10, 20, 20 units of work, resp
  - At time 20 (and 21, 22, ..., 1019) : two 1-unit jobs

SRPT:  $10 + 10 + 2000 + 1030 + 1030$

Better schedule:  $10 + 20 + 20 + 2000 + 1040$

# Multi-processor flow time scheduling (m = 1)

- SRPT (using  $m$  speed-1 processors):
  - $O(\log P)$ -competitive, where  $P$  is the ratio of the maximum job size and the minimum job size.

[Leonardi & Raz 96]

- **Lower bound:** No algorithm can be better than  $\log P$ -competitive.
- Can we do better ?

# Resource Augmentation

For many scheduling problems, an online scheduler has no way to compete with an offline scheduler (which has knowledge of the future).

A naïve approach to obtaining optimality is to allow the online scheduler to have **more resources**, such as using faster processors or more processors.

(notation speed- $x$  processors:  $x$  times faster)

Can extra resources compensate the lack of future information? How much?

# Multi-processor flow time scheduling (m = 1)

- SRPT (using  $m$  speed-1 processors):  $O(\log P)$ -competitive, where  $P$  is the ratio of the maximum job size and the minimum job size.  
[Leonardi & Raz 96]
- SRPT (using  $m$  speed-2 processors): 1-competitive  
[Phillips et al. stoc 97]
- SRPT (using  $m$  speed- $s$  processors, where  $s \geq 2$ ):  $1/s$ -competitive
  - e.g.,  $s = 4$ , SRPT is four times better than optimal offline algorithm.  
[McCullough & Torng soda 04]
- SRPT variant (using  $O(m)$  speed-1 processors): 1-competitive.  
[Chan, Lam & Liu soda 06]
- SRPT (using  $m$  speed- $(1+\epsilon)$  processors, for any  $\epsilon > 0$ ):  $4/\epsilon$ -competitive  
[Fox & Mosely soda 2011]

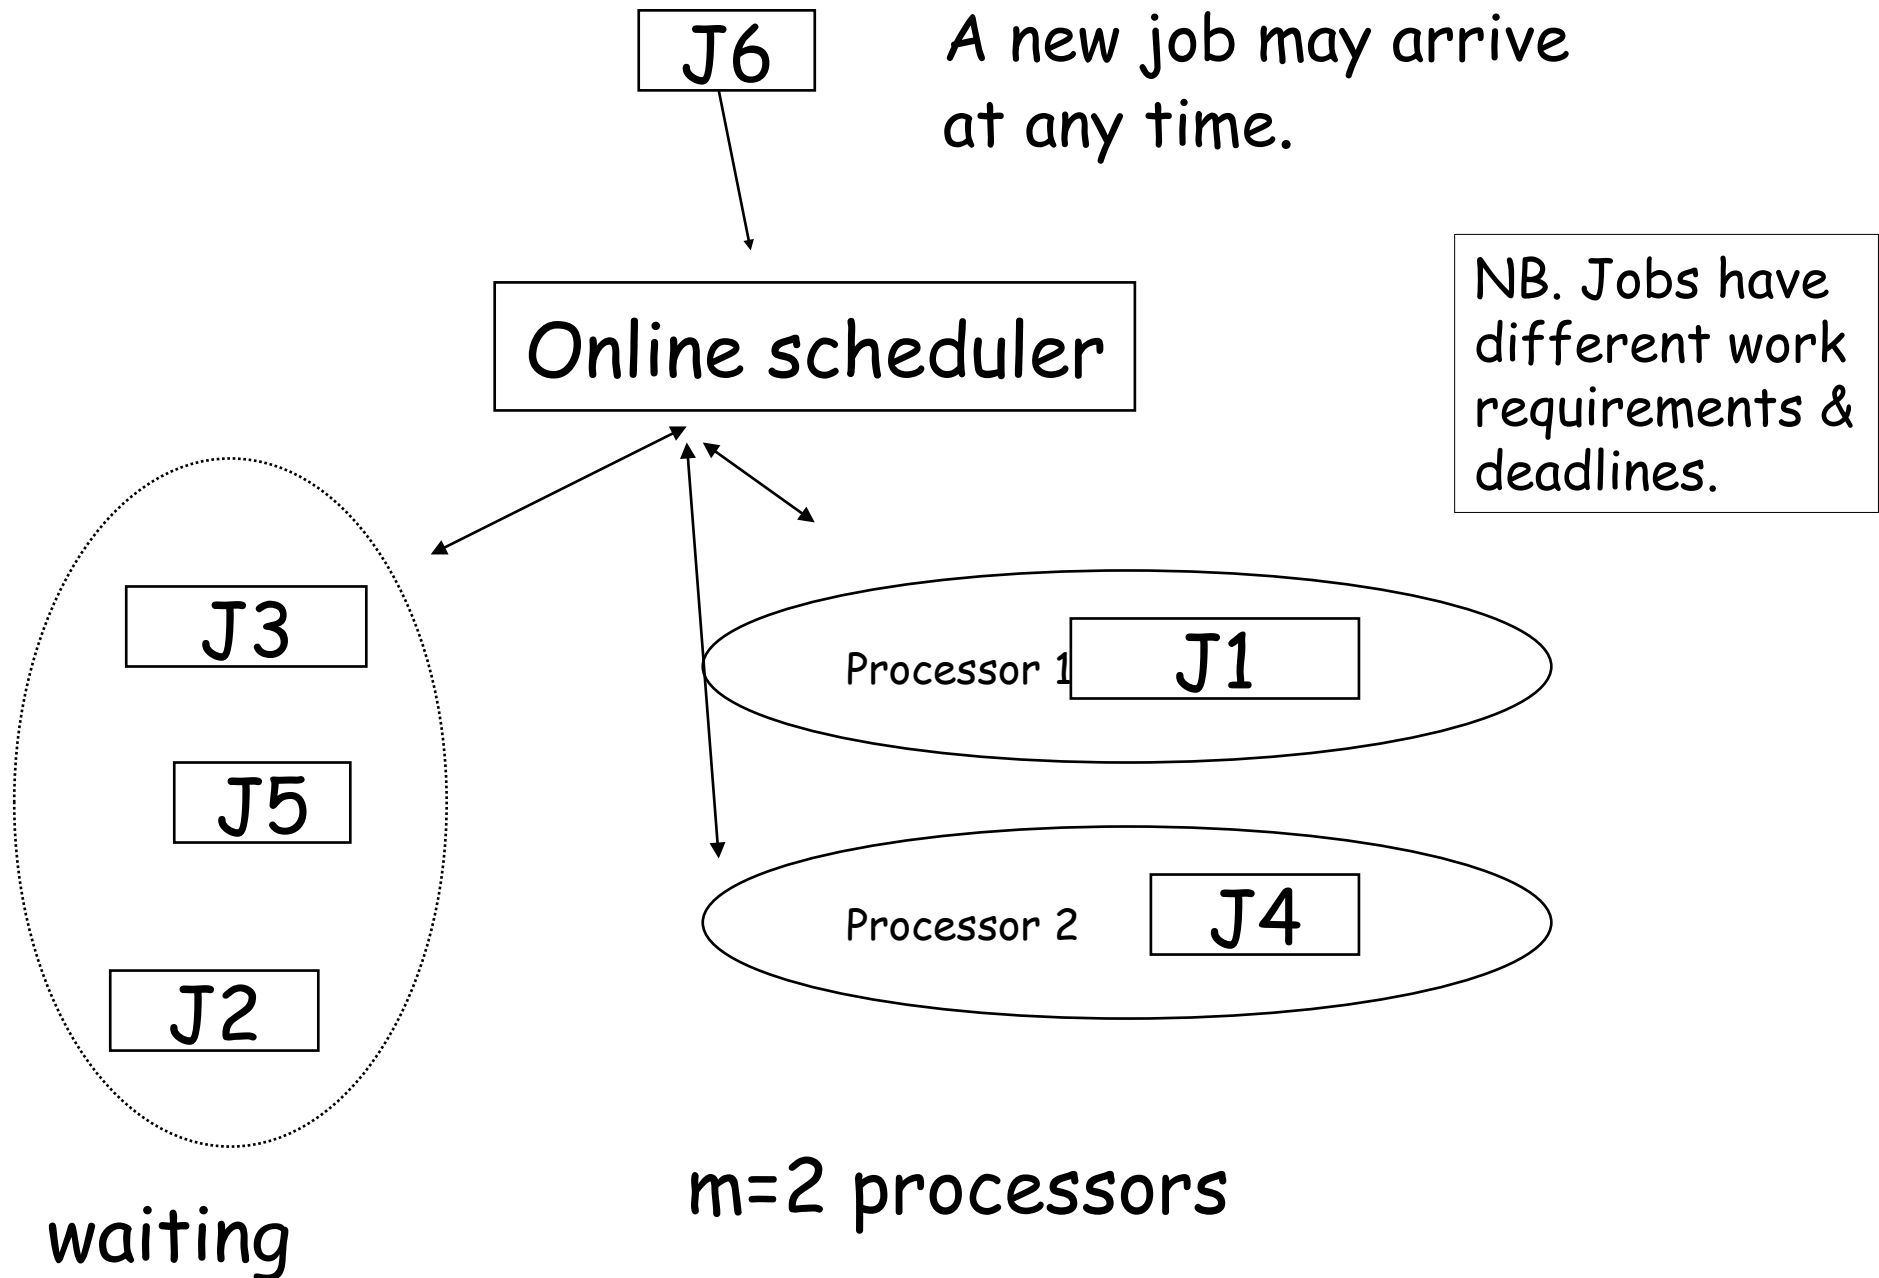
## Non-migratory online algorithm:

- IMD (using  $m$  speed-1 processors):  $O(\log P)$ -competitive [Awerbuch SPAA 02]
- IMD (using  $m$  speed- $(1+\epsilon)$  processors):  $O(1+1/\epsilon)$ -competitive [Chekuri, Khana & Zhu stoc 04]

# Online deadline scheduling

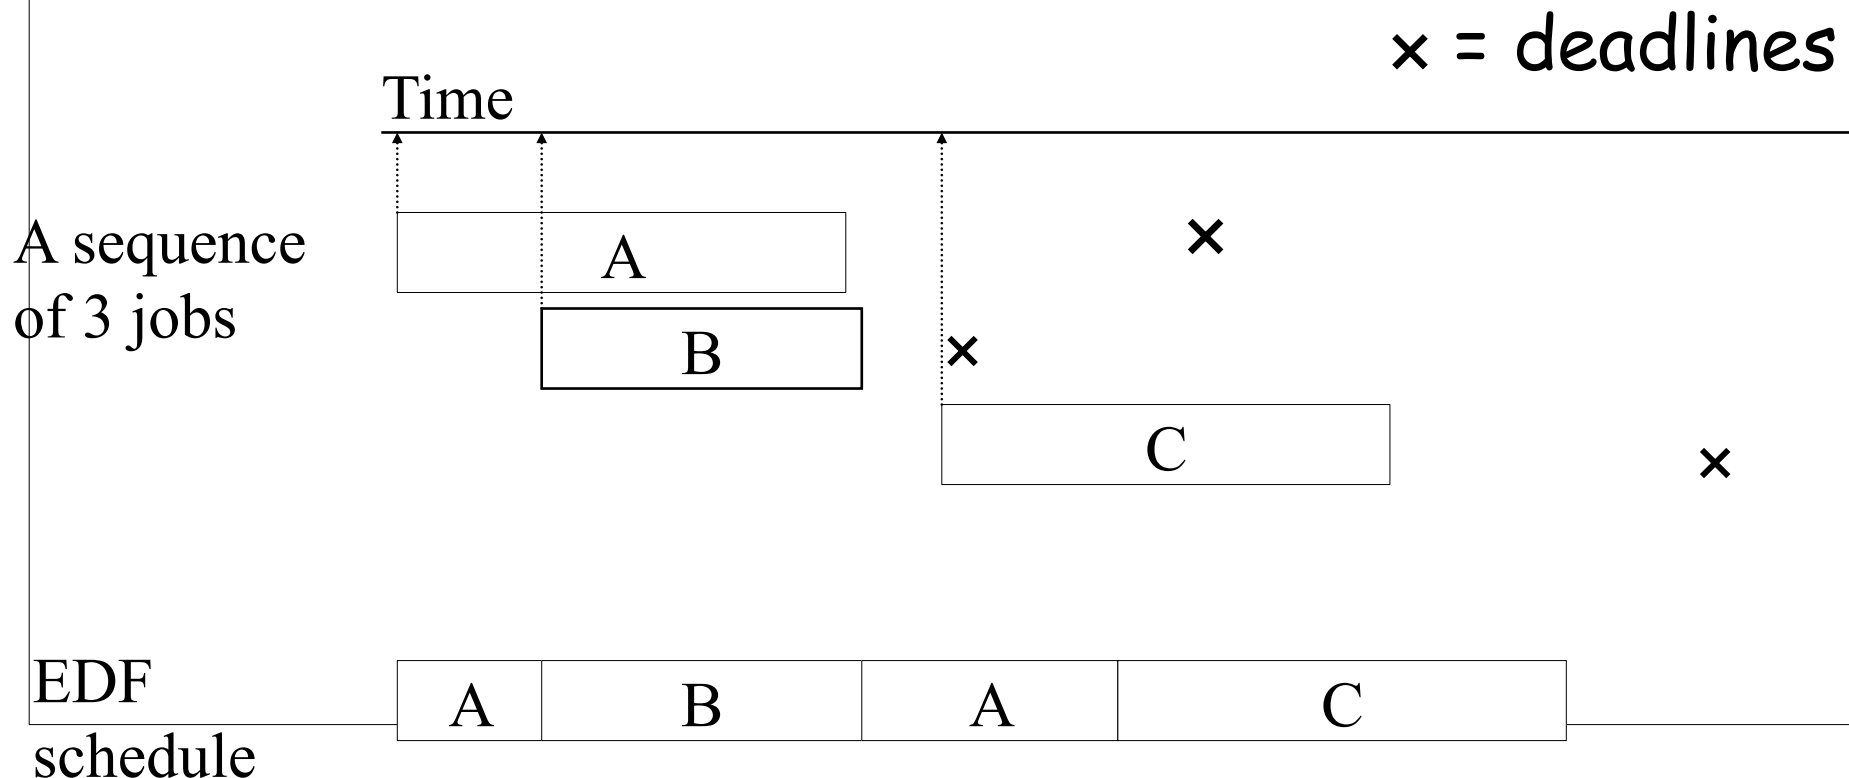
Next, we consider scheduling jobs with deadlines.

- There is a pool of  $m \geq 1$  identical processors.
- Jobs are released at arbitrary times.
- The work (processing time) & deadline of a job are known at its release time.
- An online scheduler, based on the jobs released so far, adjusts the (future) ordering of job processing.
- The aim is to complete the jobs on or before their deadlines.



# Example: $m = 1$ processor

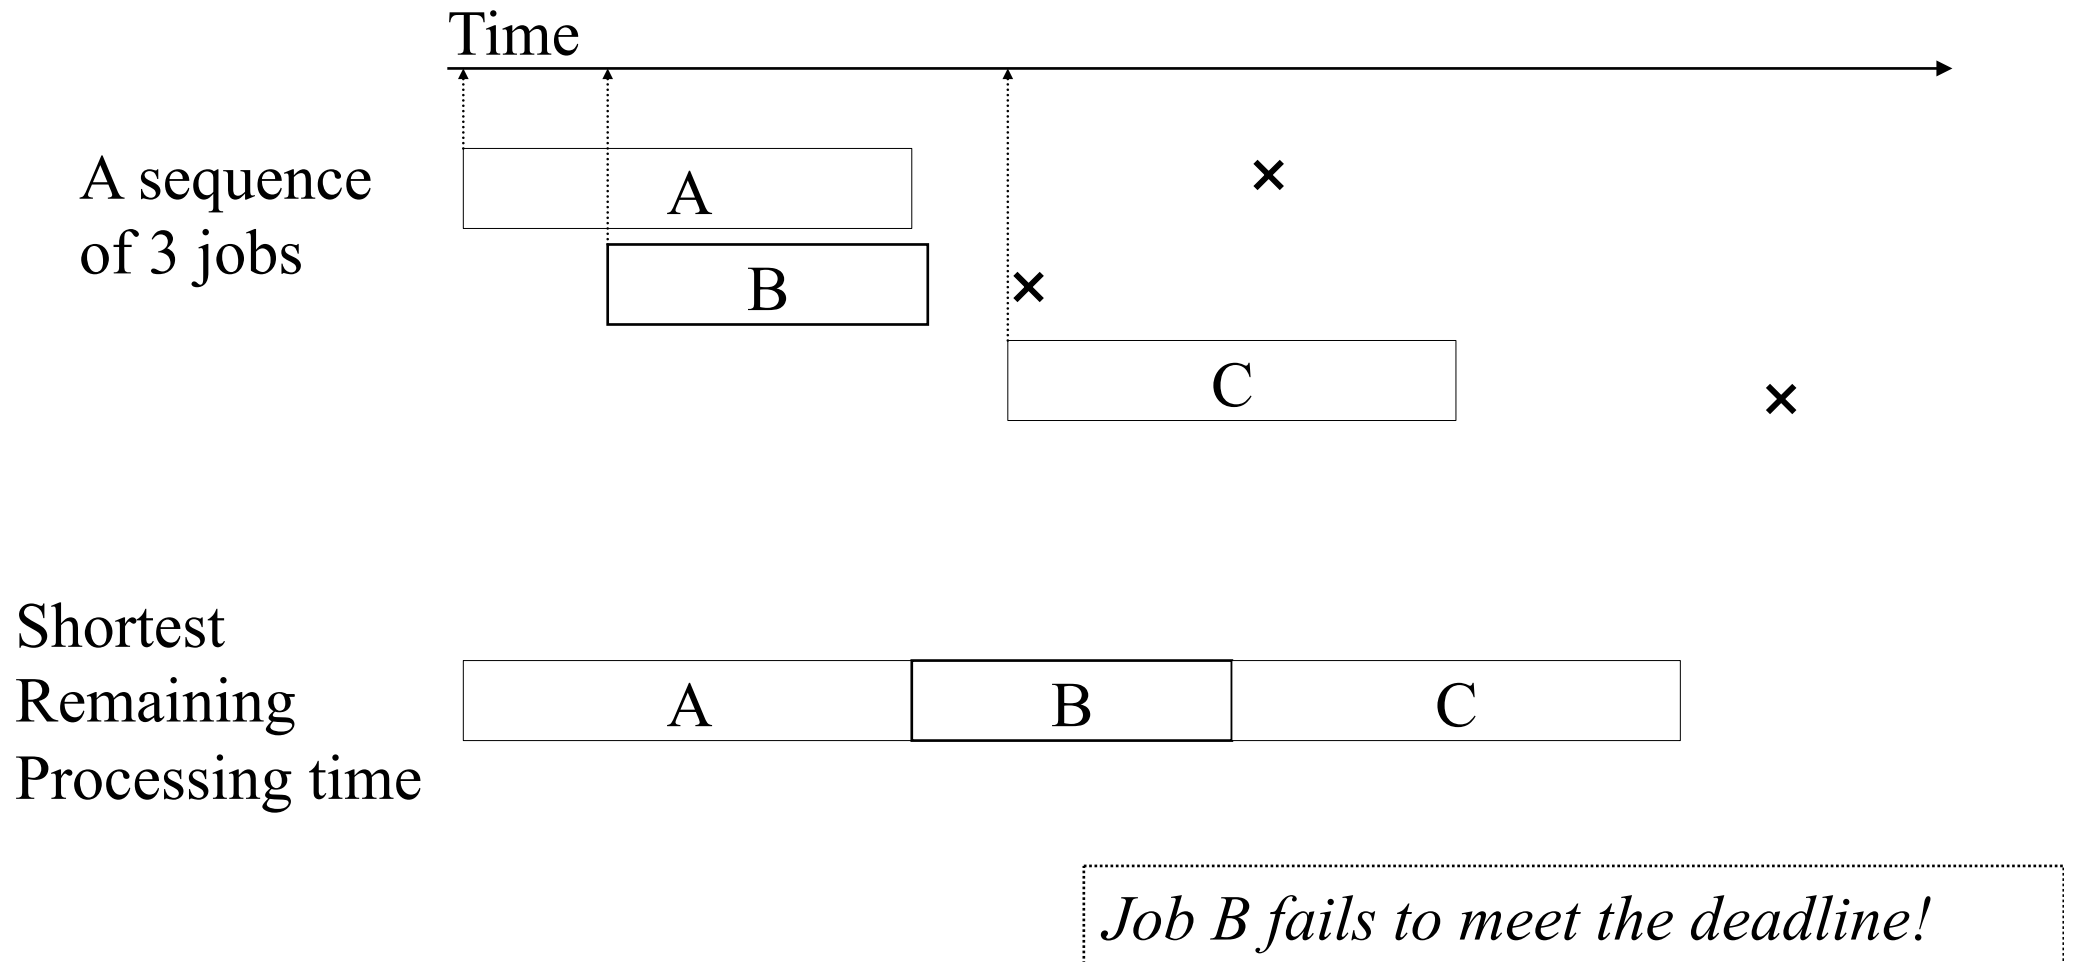
**EDF** (Earliest Deadline First) is a widely used algorithm in many real-time systems. (Ref: Deadline scheduling for real-time systems: EDF and related algorithms, Stankovic, et al., Kluwer Academic Publishers, 1988)



NB. Preemption is free.



# Example



# Hard deadline systems

- Missing the deadline of any job can't be tolerated.  
E.g., the controller of a nuclear reactor.
- An online algorithm is said to be *optimal* for *hard deadline* systems if it can schedule all jobs to meet the deadlines whenever some offline algorithm can do so.
- In other words, we are only interested in scheduling for the *underloaded* setting, i.e., when it is possible to meet the deadlines of all jobs.
- E.g., When  $m=1$ , EDF is *optimal*.

# Precisely, EDF works as follows.

- There are  $m \geq 1$  identical processors.
- Whenever a job is released or completed, check how many jobs haven't been completed.
- At most  $m$  such jobs: assign each of them to a processor.
- Otherwise, pick the  $m$  jobs with the earliest deadlines and assign each of them to a processor.
  - We assume that a processor, whenever given a job, will process it even if the deadline has already passed.
- Assumption: Every job has a unique deadline (break ties using job Ids).

# Theorem. EDF is optimal when $m = 1$

## Observation

EDF is a busy (or greedy) algorithm. I.e., EDF never lets a processor idle if there is a job not yet completed.

## Notation

Consider any input job sequence  $L$ . Suppose that an offline algorithm  $\text{Opt}$  can meet all deadlines of  $L$ .

At any time  $t$ , let

- $E(L,t)$  = the amount of work EDF has scheduled for jobs in  $L$  up to time  $t$ .
- $O(L,t)$  = the total amount of work  $\text{Opt}$  has scheduled for jobs in  $L$  up to time  $t$ .

Lemma 1: For any time  $t$ ,  $E(L,t) \geq O(L,t)$

Prove by contradiction.

Suppose that  $t$  is the first time when  $E(L, t) < O(L, t)$ .

Is it possible that at time  $t$ , EDF has processed more work than Opt on one particular job?

Is it possible that at time  $t$ , EDF has processed more work than Opt on all jobs?

Lemma 1: For any time  $t$ ,  $E(L,t) \geq O(L,t)$

Prove by contradiction.


Suppose that  $t$  is the first time when  $E(L, t) < O(L, t)$ .

Then there exists at least one job, say,  $J$ , for which EDF schedules less work than Opt up to time  $t$ .

Note that at  $t$ ,  $J$  is not yet completed by EDF.

Let  $t'$  be the release time of  $J$ . Note that  $t' < t$ .

Since EDF is a busy algorithm, EDF doesn't have any idling time during the interval  $[t', t]$ .



On the other hand, at time  $t'$ ,  $E(L, t') \geq O(L, t')$ .

And  $E(L, t) < O(L, t)$ .

During  $[t', t]$ , EDF schedules less work than Opt does.

If EDF is not idle throughout  $[t', t]$ , then EDF has scheduled the maximum possible work and can't be outperformed by Opt.

Thus, EDF must be idle at some time during  $[t', t]$ .

A contradiction occurs.


		J		
<div>EDF schedules at least as much as Opt.</div>	<div>EDF schedules less than Opt.</div>			
	J		J	



# Lemma 1 $E(L, t) \geq O(L, t) \Rightarrow$ EDF is optimal


- Suppose, for the sake of contradiction, that there exists a sequence  $I$  of jobs that
  - can be completed by their deadlines using the optimal offline algorithm  $Opt$
  - but can't be completed ... using EDF.
- Among all jobs in  $I$  that EDF fails to meet the deadline, we let  $J$  denote the one with the earliest deadline.

$d(j') \checkmark$     $d(j'') \checkmark$     $d(j''') \checkmark$     $d(j) \times$



- Let  $I' = \{ \text{Jobs in } I \text{ with deadlines} \leq \text{deadline of } J \}$ .  
(NB. Jobs have distinct deadlines; use job id to break ties)
  - With respect to  $I'$ ,  $J$  has the latest deadline.
  - $Opt$  can also meet all the deadlines of  $I'$ .

- With respect to EDF, the schedule of a particular job is not affected by the presence of jobs with later deadlines.
- The EDF schedule of  $I'$  is identical to that of the corresponding jobs in the schedule of  $I$ .
- I.e., EDF when given  $I'$  will complete all jobs except  $J$  by their deadlines.



- Let  $d(J)$  be the deadline of  $J$ . As Opt can meet the deadlines of  $I'$ ,  $O(I', d(J)) =$  the total work of  $I'$ .
- By Lemma 1,  $E(I', d(J)) \geq O(I', d(J))$ . Thus, EDF at time  $d(J)$  must have completed all jobs in  $I'$ , including  $J$ .
- A contradiction occurs.


# Lemma 1 $E(L, t) \geq O(L, t) \rightarrow$ EDF is optimal

- Use induction instead of proof by contradiction.
- Induction on job deadlines.
- Let  $J_i$  denote the one with the  $i$ -th earliest deadline.


Let  $I' = \{J_1, J_2, \dots, J_i\}$ . (NB. jobs are assumed to have distinct deadlines)

- With respect to  $I'$ ,  $J$  has the latest deadline.
- Opt can also meet all the deadlines of  $I'$ .
- By induction hypothesis, EDF on  $I' - \{J_i\}$  can meet all the deadlines.
- EDF on  $I'$  and EDF on  $I' - \{J_i\}$  have identical schedule for all jobs in  $I' - \{J_i\}$ .
- EDF on  $I'$  can meet all the deadlines of  $I' - \{J_i\}$ .
- At time  $t = \text{deadline}(J_1)$ ,  $E(I', t) \geq O(I', t)$ , and Opt meets all the deadlines of  $I'$ .
- Thus, EDF on  $I'$  must have finished  $J_1$  at time  $t$ .


- Question: When  $m = 2$ , is EDF optimal?




EDF using  
2 processors




- Question: When  $m = 2$ , is EDF optimal?



EDF using  
2 processors



Optimal  
schedule using  
2 processors



~~NB.~~ No online algorithm is optimal [Dertouzos & Mok 1989].

# Resource Augmentation

EDF using speed-2 processors.

- I.e., online:  $m$  speed-2 processors; offline:  $m$  speed-1 processors.
- A trivial result: When  $m = 2$ , EDF using speed-2 processors is optimal for **hard deadline** scheduling.
- A surprising result: For all  $m > 2$ , EDF using speed-2 processors is optimal for **hard deadline** scheduling.

[Phillips et al. STOC 97]

## Theorem 2: For $m \geq 2$ , EDF is speed-2 optimal.

- Precisely, for any job sequence  $I$ , if any offline algorithm using  $m$  speed-1 processors can meet all deadlines of  $I$ , then EDF using  $m$  speed- $s$  processors, where  $s = (1 + \frac{m-1}{m})$ , can always do so.
- When  $m = 2$ ,  $(1 + \frac{m-1}{m}) = 1.5$   
When  $m = 3$ ,  $(1 + \frac{m-1}{m}) = 1.66\dots$   
When  $m = 4$ ,  $(1 + \frac{m-1}{m}) = 1.75$   
For all  $m$ ,  $(1 + \frac{m-1}{m}) < 2$ .

## Observation

EDF is a busy (or greedy) algorithm. I.e., EDF never lets a processor idle if there are  $m$  or more jobs not yet completed. In other words, whenever a processor is idle, there are at most  $m - 1$  jobs not yet completed and each of them is currently processed by a processor.

## Notation

Consider any job set  $L$ . Suppose that an offline algorithm  $Opt$  using  $m$  speed-1 processors can meet all deadlines of  $L$ .

At any time  $t$ , let  $E_s(L, t)$  and  $O(L, t)$  denote respectively the total amount of **work** EDF (using  $m$  speed- $s$  processors) and  $Opt$  (using  $m$  speed-1 processors) have scheduled for jobs in  $L$  so far.



## Lemma 2: For any time $t$ , $E_s(L, t) \geq O(L, t)$

Prove by contradiction.

Suppose that  $t$  is the first time when  $E_s(L, t) < O(L, t)$ .  
There exists a job  $J$  for which EDF, up to  $t$ , has scheduled less work than Opt has.

Note that at  $t$ ,  $J$  is not yet completed by EDF.

Let  $t' < t$  be the release time of  $J$ .

Consider the interval  $[t', t]$ .

- The presence of  $J$  guarantees that EDF would keep at least one processor busy at any time.
- Let  $x$  be the amount of time when EDF schedules all  $m$  processors, and let  $y$  be the remaining time. I.e.,  $t - t' = x + y$ .

## Within the interval $[t', t]$

Consider job  $J$ .

- How much work EDF has scheduled  $J$ ?  
 $\geq s y$ . Reason: Whenever EDF lets a processor idle,  $J$  must be scheduled on a processor.
- Work scheduled by Opt for  $J$  is  $\leq x + y$
- Recall that EDF schedules less work for  $J$  than Opt does. Thus,  $x + y > s y$ .
- Intuitively, as  $s > 1$ , " $x + y > s y$ " favors a big  $x$ , or equivalently, a small  $y$ .

## Within the interval $[t', t]$

By definition of  $t$ ,  $E_s(L, t) < O(L, t)$ , and  $E_s(L, t') \geq O(L, t')$ .

Thus, the total amount of work EDF has scheduled during  $[t', t]$ , denoted  $\alpha$ , is strictly less than the total amount of work Opt has scheduled during  $[t', t]$ , denoted  $\lambda$ .

- $\alpha \geq s(m x + y)$
- $\lambda \leq m(x + y)$
- Therefore,  $m(x + y) > s(m x + y)$

Intuitively, the above inequality favors a small  $x$ .

# Contradiction

$$\text{I.} \quad x + y > s y$$

$$\text{II.} \quad m(x + y) > s(m x + y)$$

$$(m-1) \text{ I} : (m-1)(x + y) > s(m-1)y$$

$$(m-1) \text{ I} + \text{II} : (m + m - 1)(x + y) > s m(x + y)$$

Recall that  $s \geq 1 + \frac{m-1}{m}$ .

Thus,  $s m \geq m + m - 1$ .

A contradiction occurs!

**Lemma 2**  $\Rightarrow$  EDF is speed- $(1 + m^{-1}/m)$  optimal.

- Suppose, for the sake of contradiction, that EDF is not speed- $(1 + m^{-1}/m)$  optimal.
- I.e., there exists a sequence  $I$  of jobs that can be completed by their deadlines using Opt but not using EDF.
- Let  $J \in I$  be the job with the earliest deadline that EDF fails to meet its deadline.
- Let  $I' = \{ \text{Jobs in } I \text{ with deadlines} \leq J\text{'s deadline} \}$ .  
(NB. Recall that jobs are assumed to have unique deadlines)
  - With respect to  $I'$ ,  $J$  has the latest deadline.
  - $(I - I')$  contains jobs with deadlines  $> J$ 's deadline.
  - Opt can also meet the deadlines of  $I'$ .

- With respect to EDF, the schedule of a particular job is not affected by the presence of jobs with later deadlines.
- Thus, the schedule of  $I'$  is identical to the corresponding jobs in the **schedule** of  $I$ .
- I.e., EDF when given  $I'$  will again fail to complete  $J$  by its deadline.
- Let  $t$  be the deadline of  $J$ . As Opt can meet the deadlines of  $I'$ ,  $O(I', t)$  = the total work of  $I'$ .
- By Lemma 2,  $E(I', t) \geq O(I', t)$ . Thus, EDF at time  $t$  must have completed all jobs in  $I'$ , including  $J$ .
- A contradiction occurs.

# Trade-offs between processors & speed

- Can we reduce the speed requirement for optimality by allowing the online scheduler to use more than  $m$  processors? (technology versus money)

# Trade-offs between processors & speed

- Can we reduce the speed requirement for optimality by allowing the online scheduler to use more than  $m$  processors? (technology versus money)
- Yes.
- Precisely, for any job sequence  $I$ , if any offline scheduling algorithm using  $m$  speed-1 processors can meet all deadlines of  $I$ , then EDF using  $m + \varepsilon$  speed- $s$  processors, where  $s = (1 + m^{-1}/m + \varepsilon)$ , can always do so.
- Examples:
  - When  $m = 2$ ,  $\varepsilon = 1$   $(1 + m^{-1}/m + \varepsilon) = 1.3333$
  - When  $m = 4$ ,  $\varepsilon = 2$   $(1 + m^{-1}/m + \varepsilon) = 1.555$
  - When  $\varepsilon = 9m$ ,  $(1 + m^{-1}/m + \varepsilon) \approx 1.1$
- By choosing sufficiently large  $\varepsilon$ , we can make the speed requirement arbitrary close to 1.



# Proof Sketch

$$\text{I.} \quad x + y > s y$$

$$\text{II.} \quad m(x + y) > s[(m + \varepsilon)x + y]$$

$$(m + \varepsilon - 1) \text{ I} : (m + \varepsilon - 1)(x + y) > s(m + \varepsilon - 1)y$$

$$(m - 1) \text{ I} + \text{II} : (m + m + \varepsilon - 1)(x + y) > s(m + \varepsilon)(x + y)$$

$$\text{Thus, } s < (m + m + \varepsilon - 1) / (m + \varepsilon) = \mathbf{1 + m^{-1} / m + \varepsilon}$$

# Open problem

- Does there exist an online algorithm that uses  $m + O(m)$  speed-1 processors and can meet the deadlines of all jobs whenever it is feasible in the offline sense?

# No job migration among the processors

- EDF requires migration, i.e., a preempted job may resume execution in another processor with any cost.
- In reality, migration requires overhead and it is not realistic to expect that migration costs nothing.
- Let  $L$  be a job sequence that can be completed by an offline algorithm using  $m$  speed-1 processors
- The  $L$  can also be completed by
  - migratory online algorithms (EDF and FR) using speed-2 processors
  - non-migratory online algorithm PARK [Chan, Lam & To, SODA 2004] using speed-5.828 processors

# Overloaded systems

- *Under loaded* systems (hard deadline): there always exists a schedule that can meet the deadlines of all jobs.
- *Overloaded* systems (firm/soft deadline): there might be too many jobs and no (online/offline) algorithms can meet all the deadlines.
- In the overloaded case, an online algorithm  $A$  is said to be
  - **optimal** (work-optimal) if  $A$  always matches the optimal offline algorithm regarding the total work of jobs that are completed by the deadlines.
  - **$c$ -competitive** if the total work completed by  $A$  is always at least a fraction of  $1/c$  of that of the optimal offline algorithm.

(NB. We do not require the same set of jobs to be completed.)

# Deadline scheduling under overload

Question: In the overloaded setting, is EDF optimal for scheduling a processor ( $m=1$ )?

No.

In fact, no algorithm is optimal.

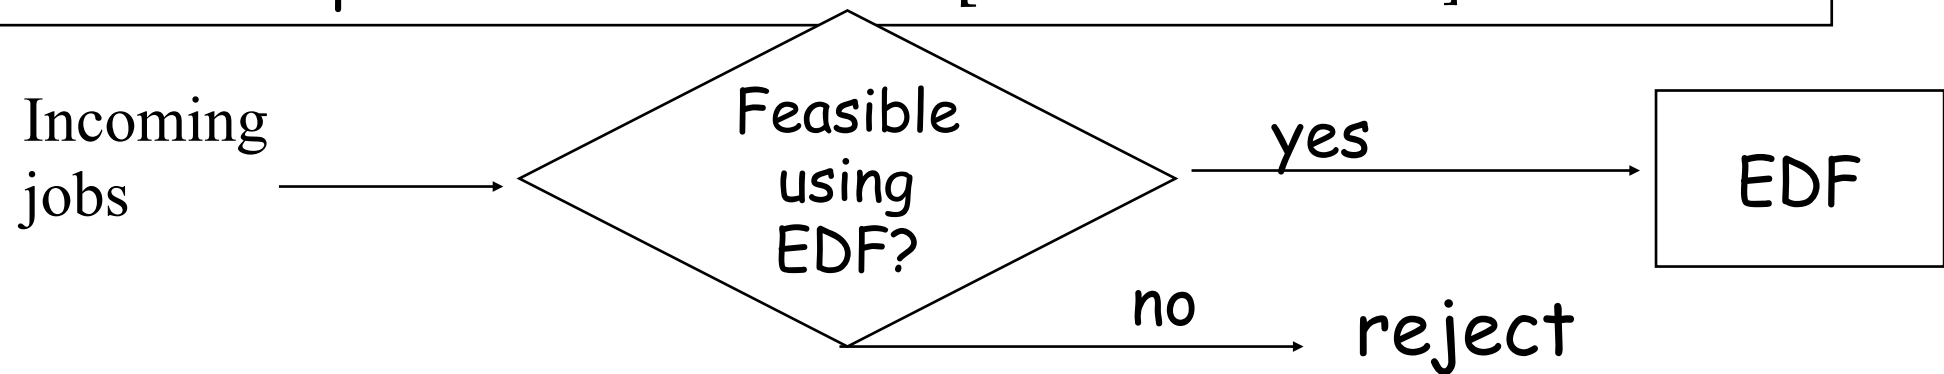
Resource augmentation: Is EDF, using a speed-2 processor, optimal for scheduling on a processor?

No.

# Online algorithms for overloaded scheduling

For scheduling one processor ( $m = 1$ ):


- $D^{\text{over}}$ : (using a speed-1 processor) is 4-competitive; there is also a lower result that no algorithm can be better than 4 competitive. [Koren & Shasha sicomp 96]
- EDF with admission control: (using a speed-2 processor) is 1-competitive. [Lam & To SODA'01]



**Lower bound:** No algorithm is speed- $s$  optimal if  $s < 1.618$  (i.e., the golden ratio.)

Suppose, for the sake of contradiction that, there is an online algorithm  $A$ , using a speed- $s$  processor, is optimal, where  $1 < s < \phi = (1 + \sqrt{5})/2 \approx 1.618$ .

Then we can obtain a contradiction with a job sequence of 4 jobs.



Offline algorithm (using a speed-1 processor) can complete both  $J_1$  and  $J_4$ .

One can prove that  $A$  will complete either  $J_3$  or  $J_4$ .